

Numerical methods project

Anders A. Søndergaard

July 5, 2011

Abstract

In this report a method is described for compressing sound files using the Fast Fourier Transform. An algorithm is implemented in C, and a few example sound files are compressed to a little less than half their size. Some theory for further improvements and techniques will be discussed.

1 Theory

A typical raw sound file consists of a series of samples, that, when fed through a digital-analog converter at a given rate, will give a sound signal that can be applied to a loudspeaker.

These samples can be thought of as a list of (t,x) pairs where t is time and x is position. If we take the Fourier transform of this list, we get a list of amplitudes and frequencies (the Fourier coefficients, c_k).

This can be exploited to compress the sound file. The human ear can hear frequencies in the range of roughly 60-17,000 Hz, so if we cut away all frequency samples not in this range, we save space. Furthermore, if the amplitude of a given frequency is very low compared to the other samples, we can ignore it too.

This is an example of results from a field called psychoacoustics. Other results are that the human ear masks certain frequencies when others are present, and perceive nonexistent frequency components under certain circumstances. For example the ear will not be able to distinguish a harmonic series where the fundamental node is missing from one where it is present.

In this small project, we only throw away frequencies out of the 60-17000 Hz range and frequencies of low intensity, but in the MP3 compression scheme, as many results as possible from psychoacoustics are used.

Normally, one splits the file into time windows of some length. There are a number of advantages of doing this in MP3 for example. It is easier to stream the file, and you won't need to have all the raw samples in memory. But more importantly, in MP3, the sample waveform can be more difficult to encode some places than others. The individual frames will then not be affected by local difficulties.

You also get a slight speedup, as you then go from a time complexity of $n \log(n)$ to $n \log(w)$, where w is the number of samples in each frame and n is the number of samples in the entire file.

When you split the file in frames and throw away Fourier coefficients, some edge effects are bound to occur, because we actually take a Fourier transform of a function folded with a square window. To counter this, MP3 overlaps frames and use other windowing functions to smoothe out the frame transitions. In this time constrained project, we don't bother, and accept some clicking in the decompressed files.

To determine the frame size we're going to use (number of samples per frame), we consider the frequency resolution needed. We (arbitrarily) set the required resolution to 10 Hz (to be able to distinguish 60 from 70 Hz). At least we need to be able to cut at 17000 Hz, and not 16000 Hz. The frame size might be optimized to allow for a greater number of samples to be thrown away, but I have not investigated this matter further. MP3 uses a frame size of a few hundred frequency samples, but spread out data for each position sample over multiple frames (not just the nearest neighbours).

When we choose 10 Hz resolution, we find that, since the number of samples $N = \text{Sample Rate} \times \text{time}$, and the k 'th fourier coefficient corresponds to the frequency $f_k = k/2N$, that we need a window size of at least 4410 samples if the wave file is encoded with 44100 samples per second, as is the case in most audio CD's. For the fast Fourier transform to be efficient, we need some power p that makes $2^p = N$. We therefore choose $p = 13$, giving us a window size of 8192 samples. It should be noted that other Fast Fourier Transform algorithms have been devised, that does not have this requirement. Our Cooley-Tukey algorithm, for example, can also be used to split the data into 3 or 4 or 5 and so on lower transforms, not just 2.

We throw away all (now complex) samples above $k = 1700$, and end up with only needing to store about 58% fewer numbers, reducing our file to 42% of its original size.

To improve on this number, we also throw away fourier coefficients that have a too low amplitude. A clever algorithm could be found that determines what too low means, but in our project we say that a minimal amplitude of 0.01 is required to be needed. This value is experimentally found to be the highest value that still gives a reasonable sound quality on the authors cheap laptop loudspeakers.

When we throw away samples like this, we need a bitmap of which samples are kept and thrown away. For the 1700 samples, this would take up 212 bytes, equivalent to 53 complex samples encoded with 16 bits of accuracy. We then need to throw away at least 3% of the remaining samples for this method to be worthwhile, which does not seem unreasonable.

On a side note, this method alone applied to an uncompressed noisy signal, could blank out the noise. If you know the frequency of the noise, you can also simply remove the offending frequency from the spectrum.

The Canonical WAVE file format

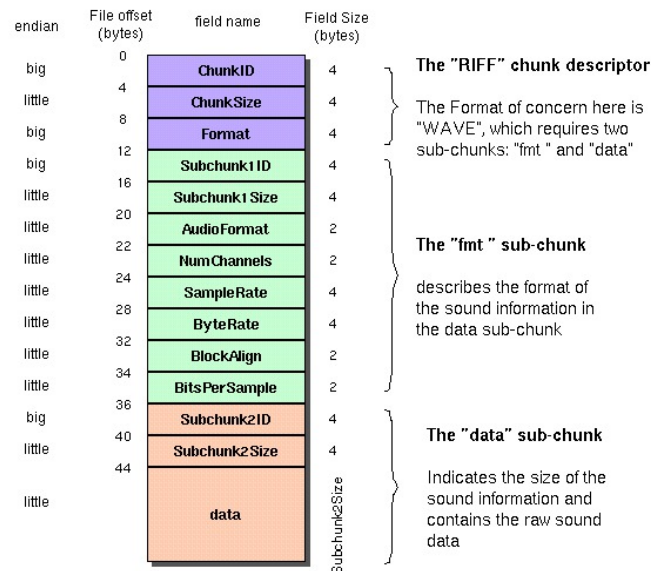


Figure 1: Canonical WAVE file format [MS WAV]

2 WAV files

A WAV file is a RIFF file (binary file of chunks) that consists of a header describing the data and the data itself placed in the “fmt ” and “data” subchunks, respectively. The samples are stored from start to end with each channel (mono, stereo, ...) interleaved so the file can be read sequentially and still produce samples for each channel. If there are 8 bits per sample, an offset of 127 is added, otherwise the samples are stored as regular 2's complement signed integers.

To read these files, a simple library was written. It provides functions for reading already existing WAV files, and functions for creating and writing to WAV files created by the program (ie. only canonical WAV files, not wav files with other headers).

3 Algorithms

For compressing WAV files, the compressor enters this main loop

```
while ((ReadSamples = readwav(...)) != 0) {
```

```
...
```

```

rmax = imax = 0; // Frame abs maximum real , imag part
for (n = 0; n < NumChannels; n++) {
    if (ReadSamples < NumSamples) { // Fill the rest with 0's
        for (k = ReadSamples; k < NumSamples; k++)
            channels[n][k] = 0;
    }
    // Prepare for FFT
    for (k = 0; k < NumSamples; k++) {
        channelspreFFT[n][k] = (complex) channels[n][k];
    }

    // FFT:
    fft(NumSamples, channelspreFFT[n], channelspostFFT[n], -1);

    // Readout:
    for (k = minSample; k < maxSample; k++) {
        // Check if we can throw away sample
        if (cabs(channelspostFFT[n][k]) > MIN_AMPLITUDE) {
            bitmap_handler_set(&header, k-minSample, n);

            channels[n][2*(k-minSample)] = creal(channelspostFFT[n][k]);
            channels[n][2*(k-minSample) + 1] = cimag(channelspostFFT[n][k]);

            // Check is a new maximum has been found
            if (rmax < fabs(channels[n][2*(k-minSample)]))
                rmax = fabs(channels[n][2*(k-minSample)]);
            if (imax < fabs(channels[n][2*(k-minSample) + 1]))
                imax = fabs(channels[n][2*(k-minSample) + 1]);
        } else {
            bitmap_handler_clr(&header, k-minSample, n);
        }
    }
}
header.real_max = rmax;
header.imag_max = imax;
if (ffwritenext(ff, &header, channels, 2*(maxSample-minSample)) != 0) {
    ...
}
i += ReadSamples;
...
}

```

What basically happens is that for each iteration in the outer loop, a frame is read from the WAVE file, a Fourier transform is taken and we store from a minimum sample (60 Hz) to a maximum sample (17000 Hz) the real and imaginary parts of the transform in another RIFF file. It is also checked if the sample is big enough to be accepted, and the proper bit in a bitmap is set or cleared accordingly.

The samples are then written with real part and imaginary part as if they were samples in another RIFF WAVE file, except each frame is prepended with a header containing a magic byte (value 42), the max imag and real part (since the samples are divided by this number to fit between -1 and 1) and the bitmap. The fourier transform preserves norm, so a waveform restricted to between -1 and 1 can still cause Fourier coefficients of up to the number of samples (8192) in size.

The file extension on the compressed files are called .ff (instead of .wav).

Decompression is almost exactly the opposite:

```
while ((ReadSamples = freadnext(...)) != 0) {
    ...

    for (n = 0; n < NumChannels; n++) {
        for (k = 0; k <= NumSamples/2; k++) {
            if (k < minSample || k >= maxSample) {
                channelspreFFT[n][k] = 0;
            } else {
                channelspreFFT[n][k] = channels[n][2*(k-minSample)] + I *
                channels[n][2*(k-minSample) + 1];
            }
        }
        for (k = NumSamples/2+1; k < NumSamples; k++) {
            channelspreFFT[n][k] = conj(channelspreFFT[n][NumSamples-k]);
        }

        fft(NumSamples, channelspreFFT[n], channelspostFFT[n], 1);

        for (k = 0; k < NumSamples; k++) {
            channels[n][k] = creal(channelspostFFT[n][k]);
        }
    }
    if (writewav(wav, channels, i, NumSamples) != 0){
        ...
    }
    i += NumSamples;
}
```

File	Original [MB]	Compressed [MB]	Ratio
Pink Floyd - Run Like Hell	45	14	69 %
Tom's Diner	28	3.6	87 %
Tool - Ænima	68	25	63 %
Sash - Equador*	35	15	57 %

Table 1: Results of compressing different files with the described algorithm. The last song was decoded from an existing MP3 file, so a lot of frequency information has already been removed.

Here the next frame is read in, and the Fourier coefficients are restored. Since we are dealing with a transform of real numbers, the last half of the coefficients are conjugate of the first half.

After the reverse transform, the samples are written to a regular wave file.

4 Results

In the folder “testcases”, a number of wav files (either copied from actual audio CD's or downloaded from the internet in FLAC format, losslessly compressed except for one) have been compressed and then decompressed. I have tried to find the song “Tom's Diner” by Suzane Vega, but have only been able to locate the music. This song was used in testing the MP3 compression, apparently because it was tricky to restore properly.

A list of the compression ratios can be found in table 1.

In the decompressed files we do find a number of places where the effects of a square envelope function is quite pronounced. We hear regular clicks (about every 20 ms) because of nonmatching frame boundaries. In a real application, this is taken care of by overlapping the frames with non-square envelopes and taking the average. This hurts compression efficiency slightly.

4.1 Odd result

In the folder “odd_stuff” a file can be found that is a result of an experiment. I tried storing only 1 (decimal) digit of the Fourier coefficients to see how sensitive the (inverse) Fourier transform is to change in accuracy of the coefficients. It sounds like the original except as if it was played on an organ. The voice still sounds almost like the original.

It turns out that 2 or 3 (decimal) digits are mostly enough to restore a song.

References

[Wikipedia] Wikipedia articles about Discrete Fourier transform and the MP3 file format and Psychoacoustics + sublinks.

[Fedorov, Dmitri] Numerical Methods, Lecture Notes.

[MS WAV] <https://ccrma.stanford.edu/courses/422/projects/WaveFormat/>
a webpage that describes the canonical WAV file format.